

Agenda

- 兩個考驗
 - 流動性
 - 風格轉變
- SaaS 的文藝復興
 - 掌握AI 重塑競爭力
- 個股

Topic 1 流動性考驗

美國公共債務突破
40兆美元，財政赤
字高懸於 GDP 的
6%。



主權債務巨獸

AI 算力無底洞



預計今明兩年全美
AI 相關信貸發行量
高達 4,500億至
5,500億美元。

26 : 2500E
27 : 4200E

全球長期
資本池



期限溢價飆升 /
收益率上行

這並非單純的流動性枯竭，而是「創紀錄的企業資本開支」與「失控的政府主權債務」
在固定收益市場中正面爭奪有限資金，直接推高長期資本的定價。

能源地緣風險：期限溢價的隱性推手

美國戰略石油儲備 (SPR)



美國戰略石油儲備 (SPR) 已跌破 3 億桶大關，創 1983 年以來新低。四季度起，官方拋儲的市場緩衝作用將消耗殆盡。

霍爾木茲海峽與紅海分流



通脹的尾部風險與聯準會政策的不確定性，迫使投資者對長端美債要求更高的期限溢價 (Term Premium)。

資產配置 I：美股的定價臨界點

美股超預期盈利

預期今年標普盈利增長
17%，明年 13%



近似無風險回報 (5-6%)

30 年期美債與投資級別公司債
(IG) 提供穩定高息

觸發條件：

一旦科技巨頭因變現推遲、地緣風險或全球需求放緩，導致企業盈利增長出現「二階導數」的放緩跡象，將觸發天量資金從高風險股權向固收資產的超級大遷徙。機構配置者正處於高度觀望狀態。

利率市場下一步及短期投資含義

- 大摩預測還有兩次加息
 - 2026 年 12 月
 - 2027 年 3 月
- 高利率維持到2027年底
- 投資含義
 - 資金的限制在提高, 瘋牛轉成慢牛
 - 殖利率提高, PER降低
 - 光有成長不夠, 要超預期

2nd 考：風向改變

訓練走向推理, 改變了什麼?

The Morgan Stanley AI Guidebook: Navigating the Age of Inference

Collaborative work across 7 teams details our views on 8 key GenAI debates, why we expect investor dollars to flow toward enablers/software and where we are seeing AI adopter signal (tech, comm services, financials, healthcare). We also detail next AI debates to consider for '27 and beyond.

8 GenAI Debates as We Enter the Age of Inference: Today's collaborative guidebook across 7 Morgan Stanley teams details our views and frameworks on the myriad of GenAI debates swirling the economy including: **1)** hyperscaler capex trajectory; **2)** GW compute capacity coming online; **3)** GenAI ROIC; **4)** the \$50-\$60tn knowledge work and consumer TAMS to be digitized; **5)** early GenAI adoption signal across the economy; **6)** the importance of open-weights models; **7)** AI credit financing; and **8)** regulatory concerns and other bottlenecks to monitor. ***We hope this guidebook is helpful in detailing our expansive views and please do let us know if interested in any of our interactive models across the space.***

MORGAN STANLEY & CO. LLC

Brian Nowak, CFA

Equity Analyst

Brian.Nowak@morganstanley.com

+1 212 761-3365

Stephen C Byrd

Equity Analyst

Stephen.Byrd@morganstanley.com

+1 212 761-3865

Adam Wood

Equity Analyst

Adam.Wood@morganstanley.com

+1 212 761-3656

Joseph Moore

Equity Analyst

Joseph.Moore@morganstanley.com

+1 212 761-7516

Michelle M. Weaver, CFA

Equity Strategist

Michelle.M.Weaver@morganstanley.com

+1 212 296-5254

Ariana Salvatore

Equity Strategist

Ariana.Salvatore@morganstanley.com

Lindsay A Tyler

Credit Analyst

Lindsay.Tyler@morganstanley.com

+1 212 761-2734

Vishwanath Tirupattur

Strategist

Vishwanath.Tirupattur@morganstanley.com

+1 212 761-1043

David Arcaro, CFA

Equity Analyst

David.Arcaro@morganstanley.com

+1 212 761-1817

Julian Herrera

Research Associate

Julian.Herrera@morganstanley.com

+1 212 761-1784


Gregory Gao

Research Associate


Gren.Gao@morganstanley.com

+1 212 296-3125

Monetization in Mobile Era: Indexed Relative Stock Performance



Early Monetization in AI Cycle: Indexed Relative Stock Performance



Relative Stock
Performance During
Early AI Era:

資本支出的頂峰與放緩

超大規模雲端服務商 Capex (2025-2028)



2027 年超大規模雲端服務商 Capex 將暴增 60% 達 1.5 兆美元，但 2028 年成長率將急遽放緩至 12% (1.6 兆美元)。

資料來源：摩根士丹利研究

戰略啟示：這種支出的放緩，伴隨著雲端/軟體營收的規模化與現金流改善，是驅動機構投資者將資金從硬體邊際輪動至「賦能者與軟體層」的關鍵催化劑。強烈看好 AMZN, GOOGL, MSFT, META。

Why We See Capex Growth Slowing in '28 and Where We Could Be Wrong

Why We Think it Slows

- **Bottom-Up Chip/Rack Work:** We build our capex on a bottom-up basis based on chip and rack counts/estimated shipments and cost/chip... and there are limits on the numbers of chips that are built/available to purchase.
- **Outyear Pre Builds are Material... and Have Limits:** We estimate a substantial portion (in some cases 30-40%) of '26 capex is being spent on long-dated land, buildings, equipment, power for data centers to open in '27/'28/'29. Hyperscalers are building further out for **strategic reasons** (being able to light up compute as it comes available) and **political reasons** (getting ahead of potential further political backlash against data center builds if the Executive branch changes parties). While forward building will continue, there are limits as to how far out these pre-builds will go.
- **Physical Limits:** We see physical limits (labor, materials, power, etc.) also holding back further substantial growth off of already large numbers and rapid pace of build.
- **META '28 Downward Capex Potential:** The extent to which META chooses to leverage hyperscalers (AWS/Azure) to scale its family of models and API (as it did in the past with Llama) would lower long-term compute needs and potentially put downward pressure on '28 capex.

Where We Could Be Wrong?

- **Component Inflation:** Component inflation (memory, cost of energy, etc.) remains high and could continue to creep higher. Next generations of chips are also more expensive... though with better performance and (we believe) better revenue opportunities.
- **Larger-than-expected chip production in '29 and beyond:** Could lead to more prebuild in '27/'28.

September 7, 2026 09:32 PM GMT

Tech Bytes | **Asia Pacific**

Which AI Bottleneck To Invest In?

Astra GPT-6 strengthens the AI bottleneck narrative. The risk of under-investing now outweighs the risk of over-investing given investors' reduced positioning, macro ambiguity and a wall of worry following recent draw-downs. Ambiguity has value and we update our best global opportunities.



MORGAN STANLEY & CO. INTERNATIONAL PLC+

Shawn Kim

Equity Analyst

Shawn.Kim@morganstanley.com

+44 20 7677-1018

MORGAN STANLEY ASIA LIMITED+

Duan Liu

Equity Analyst

Duan.Liu@morganstanley.com

+852 2239-7357

MORGAN STANLEY & CO. INTERNATIONAL PLC+

Cindy Huang

Equity Analyst

Cindy.Huang@morganstanley.com

+44 20 7425-2915

MORGAN STANLEY & CO. INTERNATIONAL PLC, SEOUL BRANCH+

Ryan Kim


Equity Analyst

Ryan.G.Kim@morganstanley.com

+82 2 399-4939

資本支出(流量)增速放緩, 但算力(存量)是在疊加的

We See a Path to ~145GW of Compute Capacity Online by '28...





Inference Performance
and Efficiency Drive
Company Results


DeepSeek V4 Pro 0813 1.6T Agentic Total Tokens per \$1 TCO vs. P90 Interactivity

Cost Tier: Owning at Large Hyperscaler Volume ⓘ Updated: 2026-09-10 Source: SemiAnalysis InferenceX™

TCO \$/chip/hr: VR200: 3.61 GB300: 2.31

Source: [SemiAnalysis Market July 2026 Pricing Surveys & AI Cloud TCO Model](#) ↗

Note: This deployment is not hosted at inferencex.semianalysis.com and is not affiliated with or endorsed by SemiAnalysis. Data shown here may be unofficial, modified, or out of date — visit the official site for authoritative InferenceX™ results.



Reset filter Clear rulers (3) Quick Filters ⓘ

- ☒ Vera Rubin NVL72 (TRTLLM)
- ☒ GB300 NVL72 (Dynamo TRTLLM)
- ☐ GB300 NVL72 (Dynamo SGLang)
- ☐ GB300 NVL72 (Dynamo vLLM)
- ☐ GB200 NVL72 (Dynamo vLLM)
- ☐ B300 (SGLang)
- ☐ B300 (vLLM)
- ☐ B200 (SGLang)
- ☐ B200 (vLLM)
- ☐ MI355X (ATOM')
- ☐ MI355X (MoRI SGLang)
- ☐ MI355X (SGLang)

☒ Optimal Only ⓘ

☐ Gradient Labels


▼ Advanced

- ☐ Log Scale
- ☒ Labels
- ☐ High Contrast
- ☒ Parallelism Labels
- ☐ Line Labels
- ☒ Perf Ruler ⓘ
- ☐ # Concurrent Sessions

目前的獲利能力還是很不錯的, 3種模式如下:

We Present 3 Frameworks Detailing GenAI ROIC of 25-50%

We estimate ROIC across 3 different monetization vectors including: **1)** hyperscaler renting out GPU/compute capacity; **2)** model builders serving their API through their own, internal infrastructure; and **3)** model builders serving their API through 3P infrastructure.



NOPAT : EBIT x (1-稅率)

■ IT Depreciation ■ Non-IT Depreciation ■ Energy & Other Opex ■ Compute Rental Costs ■ Tax


可能低估了IaaS ROIC！

- Nebius Group 宣布自 2026 年 10 月 1 日起局部調升 On-demand 算力租金：
 - 其中 Nvidia H100 租用價格由每小時 US\$3.85 調升 16.9% 至 US\$4.50；
 - 次世代旗艦 B300 算力租金更大漲 21% 從每小時US\$7.85 調升到 US\$9.50。
 - 除此以外，Nebius 也調高部分AMD/Intel x86 CPU租用價格達25%，及存儲記憶體租用價格達41%，Nebius 同步確認其簽約容量全速朝向年底 5 GW 目標邁進。
- 同日，CoreWeave 官方宣布自 6 月底以來持續以更高價格簽約算力，其第三季度簽訂的 3 至 6 個月短期客戶合約，年化定價達到每 MW 約 4,000 萬美元（折合 1 GW 算力租用合約每年高達 400 億美元），合約總用電容量自 3.7 GW 增至 4.2 GW。

Gen AI 的TAM 估計

GenAI Should Be Faster for Multiple Reasons... but Even 4% Penetration in Year 2 (Conservatively Saying its '27 Here) Could Imply ~\$800bn of GenAI Enterprise Spend

- Using the public cloud adoption cycle as an analogy for future AI adoption, we see ~\$800bn of Enterprise AI spend by 2027 or about 4% penetration of the broader Knowledge Work TAM.
- However, **AI should diffuse faster than public cloud** for a few reasons including: **1)** AI doesn't require wholesale infrastructure migration; **2)** shorter time-to-value from impactful productivity gains; and **3)** digital is known and enterprises face a more substantial competitive opportunity cost if they fall behind peers on AI adoption.



利潤分配

- 基礎設施穩收租金
 - Google/MSFT/AMZN/SPCX/Neocloud
- 模型層面臨利潤擠壓與分化
 - 反：Token routing 到開放權重
 - **正：frontier model 還是供不應求**
- 應用/第一方數據平台與賦能者吃掉大部分價值與增長槓桿

大摩0721 報告

很**基本**且**全面**的改變了當下選股的邏輯

July 21, 2026 04:48 AM GMT

Software

The Moat & The Journey – Assessing Software Durability + Growth in the Age of AI

Software is evolving to address the risks + opportunities from AI, just as it has in prior cycles. To assess implications, we introduce a Moat & Journey framework and come away with conviction on MSFT, PANW, CRWD, NET, DDOG, SNOW, and SHOP and assume coverage of CRM/INTU at EW and ADBE/WDAY at UW.


NORTH AMERICA INSIGHT

MORGAN STANLEY & CO. LLC		
Adam Wood		
Equity Analyst		
Adam.Wood@morganstanley.com		+1 212 761-3656
Sanjit K Singh		
Equity Analyst		
Sanjit.Singh@morganstanley.com		+1 415 576-2060
Meta A Marshall		
Equity Analyst		
Meta.Marshall@morganstanley.com		+1 212 761-0430
Elizabeth Porter, CFA		
Equity Analyst		
Elizabeth.E.Porter@morganstanley.com		+1 212 761-3632
Josh Baer, CFA		
Equity Analyst		
Josh.Baer@morganstanley.com		+1 212 761-4223
Chris Quintero		
Equity Analyst		
Chris.Quintero@morganstanley.com		+1 212 761-1686
Jonathan Eisenson		
Research Associate		
Jonathan.Eisenson@morganstanley.com		+1 212 761-2808

機會與威脅


- 軟體產業正從過去由數位轉型與雲端運算驅動的成熟期（中位數營收成長率從 5 年前的約 20% 放緩至約 10%），**跨入全新變革週期**
- LLM 與擴散模型（Diffusion Models）能理解占企業 80% 以上的非結構化數據（郵件、簡報、PDF 等）與多媒體內容，**大幅拓展軟體能自動化的業務範疇**。
- 企業自研（DIY）、原生 AI 新創（AI-native startups）及大型語言模型供應商（LLM providers）帶來**全新競爭壓力**
- 除了傳統雲端基礎設施支出，還新增了**高昂的 AI 模型運算與 Token 成本**。
- **股價反應了威脅, 但沒看到看機會**

Exhibit 9: Software 10-Year Price Performance Vs. Market



Source: Bloomberg

Exhibit 10: Software 1-Year Price Performance Vs. Market



Source: Bloomberg

對SaaS 的誤解

- 前沿實驗室輕易取代？
- 無頭架構（Headless）失去客戶黏著度？
- 席位遭侵蝕，收費基礎下降？
- 毛利率結構性走低？
- 前沿模型token cost 價格上漲迅速？

Market View：前沿實驗室輕易取代SaaS？

MORGAN STANLEY VIEW

Labs AND select software vendors both accrue value

Raw intelligence keeps improving on academic benchmarks, but progress on real-life knowledge work has plateaued.

The "Jagged Frontier": LLMs are powerful on some tasks, fail on others. Labs stay inside the frontier — coding, general knowledge work (summarization, drafting, ideation), content generation.

App vendors win by augmenting LLMs with domain-specific context/data at the workflow layer (healthcare, legal, finance). Infra delivers data/context; cyber secures agentic workflows.


大語言模型的能力邊界並不均勻：LLM 在某些特定任務表現極為強大，但在其他任務上卻容易遭遇挫敗。

前沿實驗室（Labs）的主場：大模型實驗室（如 OpenAI、Anthropic 等）的優勢仍聚焦在能力邊界之內——包含程式編寫（Coding）、通用知識工作（文本摘要、草稿撰寫、創意發想）以及內容生成

垂直應用軟體廠商的護城河在於將 LLM 與特定領域的專業上下文及數據（Domain-specific context/data）深度結合，主導高價值的專業工作流

價值向上層遷移：過去經濟價值高度集中在應用層與使用者介面；但在 **Agent** 架構下，軟體的長期差異化正向技術棧頂層移動，關鍵轉向工作流程編排、領域專業知識以及專有資料（**proprietary data**）

Exhibit 12: More Value Accrues to the Workflow Layer in an Agentic World



Market View : 無頭架構 (Headless) 失去客戶黏著度?

MORGAN STANLEY VIEW

Open AND closed ecosystems will coexist

Agreed: APIs and agent-ready infrastructure become critical, and UI-only moats are threatened.

But not everyone goes headless. UI-moat or thin-data vendors must; vendors on regulated, high-gravity data can stay relatively closed — and will.

Result: a stack that opens at the front and consolidates into governed ecosystems at the back. Closed ecosystems have more right to own orchestration where the stack is already standardized on them.

當軟體失去前端 UI 時，會不會失去客戶心智與接觸點，最終淪為被上層 AI 或聚合器隨意切換、比價的「無名管線 (Dumb Pipe) 」？

大摩觀點：「掌握受監管、高引力數據 (Regulated, high-gravity data) 的後端廠商，不會輕易被替換，反而更有權力主導編排 (Own Orchestration) 」。

SALESFORCE — OPENS AT THE FRONT

UI-led footprint pushes toward headless — prioritizing APIs and agent access as interfaces commoditize.

SAP — CONSOLIDATES AT THE BACK

Regulated, high-gravity ERP data supports a relatively closed, governed ecosystem with a stronger claim to orchestration.

Market View：席位遭侵蝕，收費基礎下降？

MORGAN STANLEY VIEW

Unit of value changes, economics don't

Evolving pricing isn't new: vendors and customers have always negotiated around value, workflow criticality, switching costs, and bargaining power.

GenAI changes the unit of value, not software economics — owners of mission-critical workflows with measurable ROI keep pricing power.

Expect a hybrid seat-plus-consumption model for seat-based vendors; pure-consumption vendors capture AI upside sooner. Implications: new disclosures to underwrite the shift, and more cyclical as consumption mix rises.

GenAI 改變的是「價值的衡量單位」（例如從人頭席位轉向產出/使用量），但沒有顛覆軟體的底層經濟邏輯。那些掌握關鍵任務工作流（Mission-critical workflows）且能為客戶帶來具體可衡量投資報酬率（Measurable ROI）的軟體廠商，依然能保有強大的定價權。

BENEFICIARIES

Consumption-led models

SNOW, DDOG, MDB, NET, FROG, SHOP, TOST, KVYO, AMPL

Market View：AI 將使毛利率結構性走低？

MORGAN STANLEY VIEW

GM pressure, but EBIT defended — flat to up

Perception over-focuses on GM dilution and underrates the levers to defend gross profit dollars and expand EBIT.

Inference deflation, prompt caching, model routing, batching, open-weight and domain-specific models, and pricing all moderate AI COGS as products scale.

Illustrative base case: ~flat EBIT margins despite GM pressure; we lean bullish — architecture and pricing limit blended GM pressure while workflow-driven OpEx leverage expands EBIT.

Analog: software-plus-payments — lower-GM revenue can still be attractive at EBIT when it isn't R&D/S&M intensive.

毛利率下降, 但增量 AI 功能是向既有客戶以加價或用量計費形式變現, 收入基礎擴大故費用率下降更快, EBIT 反而提高

- SW cost 高的公司受益 MSFT, MSFT、SAP、ORCL、WDAY、NTNX
- 基於交易量/支付（GMV）計價的應用軟體

Market View：前沿模型token cost 價格上漲迅速？

MORGAN STANLEY VIEW

Token scrutiny is real, but not uniform

Risk is concentrated, not broad: the exposure sits where tokens fund experimentation — prototypes, agent loops, and code that never ships.


Trade-down is recalibration, not demand destruction — customers reserve frontier models for high-value, high-risk work (“no F1 car to the grocery store”).

Workflow-rich incumbents become the optimization layer: routing simple work to cheap models, reserving frontier for risk, capping low-value usage.

成熟的 SaaS 巨頭（如 **ServiceNow**、**Salesforce**、**Datadog**）直接幫企業做好任務分流（簡單分類用便宜的開源/微調模型，複雜決策才呼叫前沿大模型），企業反而更願意把錢付給能幫他們「省 Token、控預算」的軟體平台。


護城河 (The Moat)

評估當前競爭地位的持久性。



演進旅程 (The Journey)

評估適應下一代軟體週期的技術與商業能力。



頂級軟體公司不僅能在今天難以被取代（深城河），更能利用AI在下一個週期實現價值複利（短旅程）。

核心精選：兼具護城河與進化力的「頂尖企業」

在追蹤的 81 家企業中，僅少數能完美契合雙重框架。

MSFT

規模化複合增長

涵蓋雲端、數據與生產力的無處不在平台。

PANW

任務關鍵型安全

不斷擴展的 AI 威脅情報與自動化平台。

CRWD

雲端原生防禦

以 AI/ML 為核心的輕量級代理與高留存率。

SHOP

商業網路效應

具備強大生態系統與應用程式商店飛輪。

NET

邊緣基礎設施

開發者親和力與快速擴展的 AI 推理能力。

NOW

企業工作流程

深度嵌入企業且具備 AI 自動化潛力的紀錄系統。

DDOG

雲端可觀測性

強勁擴展引擎與 AI 驅動的監控自動化。

SNOW

數據引力

驅動全新 AI 工作負載的數據雲平台。

週期排序：軟體子板塊的戰略佈局矩陣

	今日 (第一階段)	中期 (第二階段)	後期 (第三階段)
基礎設施軟體	率先受惠於技術債升級與AI現代化需求	持續增長	成熟穩定
網路安全	立即防禦AI驅動的網路威脅	自動化安全營運中心 (SOC)	邊緣防護普及
應用程式軟體	面臨定價模型轉型的摩擦	核心AI能力開始自動化營運	完成商業模式轉型，實現全面變現

關鍵洞察：基礎設施與網路安全板塊的表現已超越SaaS近50%，完美反映了此週期排序。

AI 變現時序：價值流動的 3 個階段



公司介紹

MSFT

ORCL

台虹

定穎

MSFT & SPCX

MSFT短綫上缺乏繼續向上的題材(350→500)

但若宣佈向SPCX 租用算力, 兩家都將重估獲利

傲視群雄：唯一完美兼具雙優勢的企業


🏆 Moat 排名
#1 (50分)

🏆 Journey 排名
#5 (40分)

在覆蓋的 81 家公司中，僅有 5 家能同時在 Moat 與 Journey 擠進前 15 名。微軟不僅位居榜首，更是得分總和最高的「最佳運動員」。



絕對防禦：滿分 50 分的護城河溯源



生產力系統記錄 (M365)

深入全球企業日常運作，
成為無可替代的系統記錄
(System of Record)。

底層雲端與數據 (Azure)

作為企業不可或缺的數位
基建，掌握最深層的 IT
營運真相。

生態圈與轉換成本

無與倫比的平台廣度與深度整
合，創造極高昂的轉換成本，
構築堅不可摧的護城河壁壘。

贏在「先買後建 (Buy then Build)」的週期起跑點



軟體產業的 AI 週期遵循「先建置基礎設施，後發展應用程式」的規律。
微軟作為基礎設施的絕對領導者，比起純應用層 SaaS 公司，處於更有利、更早期的變現位置。

總結：無可取代的 AI 複利機

勝利的 5 大關鍵要素	微軟 (MSFT) 具備狀態
1. 無匹敵的系統記錄護城河 (System of Record Moat)	[✓]
2. 贏在「先買後建」的基礎設施優勢 (Build Advantage)	[✓]
3. 最強的 AI 應用與工作流程市占 (Copilot Context)	[✓]
4. 頂級的財務與營運槓桿效率 (Operational Leverage)	[✓]
5. 具備強大下檔保護的規模估值 (Attractive Valuation)	[✓]

Exceptional Moats. Strong Journeys. Outsized Long-Term Compounder.


微軟不僅是當下最難以撼動的防禦堡壘，更是未來 AI 征途上最銳利的矛。

On a GB300 Cluster, Revenue Per Gigawatt Climbs 8× Up the AI Stack SpaceX Prices Between Neocloud Rental and Model-Layer Tokens

Annualized revenue per gigawatt by layer — \$B per GW per year. Multiples shown vs. neocloud 5-yr IaaS average (= 1.0×).

Sources: SemiAnalysis Tokenomics Model, SemiAnalysis Inference Simulator

sales@semianalysis.com




Microsoft Azure: Y/Y Revenue Growth by Component

Source: SemiAnalysis Tokenomics Model

sales@semianalysis.com



Indicative scenario: IF Microsoft contracts 3GW from SpaceX, fully live by YE2027 at \$100M/MW/yr = \$300B exit ARR



GPT tokens on SpaceX compute are served through Azure Foundry - incremental to the base Foundry line, banded separately here. Illustrative, not a forecast.

Period-end MW: 100 (3Q26) > 400 > 800 > 1,400 > 2,100 > 3,000 (4Q27)

Amber % = Azure y/y incl. SpaceX upside Grey % = Azure y/y, base forecast ex-SpaceX

Revenue recognised on average MW live in period

SpaceX Turns Into a Compute Company in Eight Quarters

AI Compute Goes From Nothing to 77% of Run-Rate Revenue




Annualised revenue run-rate by segment — exit-month revenue x 12, \$B

2026 anchored to signed contracts · 2027 illustrative at 10 GW, with 2027 additions split 50% Grok / 50% contracted at \$50 per watt per year

Source: SemiAnalysis Tokenomics Model

sales@semianalysis.com



Orcle

算力交付進入加速


OpenAI在月初推出GPT-6 Astra，強調模型在電腦操作、軟體工程、網路安全及專業工作等方面都有大幅提升。這款模型上市後，獲得企業用戶及開發人員的熱烈反應，也促

使Anthropic的潛在投資人重新評估OpenAI是否可能開始從Anthropic手中搶走市占率。

據企業支出管理平台Ramp最新數據，GPT-6 Astra約占企業AI支出的13%，Anthropic的Claude Fable則約占8%。另外在AI模型轉接平台OpenRouter，過去一周用戶在OpenAI模型的支出高於Anthropic模型，這也是逾2年半以來OpenAI首度在該指標居於領先地位。

图 4 累计上线容量假设：从 1.2GW 到 6.0GW

以 FY2026 末 1,200MW 为期初基数，逐季叠加新增；每季期初容量=上一季度期末合计



GS Forecast

	5/26	5/27E	5/28E	5/29E
Revenue (\$ mn) New	67,358.0	89,098.5	128,926.5	174,131.5
Revenue (\$ mn) Old	67,358.0	89,013.5	125,852.8	165,807.7
EBITDA (\$ mn)	36,548.0	52,166.7	77,440.4	113,219.2
EBIT (\$ mn)	28,925.0	35,485.0	46,498.0	69,686.3
EPS (\$) New	7.66	8.07	10.27	15.59
EPS (\$) Old	7.66	7.72	10.84	15.07
P/E (X)	27.4	18.6	14.6	9.6
Dividend yield (%)	1.0	1.4	1.4	1.4
Net debt/EBITDA (X)	2.7	2.3	2.0	1.3

	8/26	11/26E	2/27E	5/27E
EPS (\$)	1.95	1.81	2.04	2.27

TP >200, 40% up

Taiflex

TPFE 研發速度比M10 快, 唯一能趕上Switch tray MP

nVidia 機櫃動向

Oberon


Phase 1 • CPO Scale-Up

機櫃間 inter-rack | 2H27 (E)



Rubin Ultra NVL576 : 8 x NVL72


Kyber



Feynman NVL 1152 : 8 x NVL144

VR300 NVL576(高配版, 10月底放榜)

- 兩層的網路架構, NVswitch 的數量翻倍, PCB 的走綫更複雜, 升級到TPFE



- 一年的量225萬片PTFE CCL
 - 20% 的機櫃採用tpfe
- 生益&other: 台虹 = 7:3
- 從Q4'27 開始出貨
- 2028每個月6萬片, 每片\$400
 - NTD7E/月
 - EPS 1/月

如果中獎TP > 400, 45% Up, 把握:中

Feyman NVL144(H4'27 定案)

- 型態：
 - 單一機櫃內部採用 Kyber 正交垂直盲插架構，直接容納 144 顆 GPU。
 - 100% 機櫃採用 tpfe
- 傳輸：
 - 走極限「純銅互連」，藉由 Kyber 縮短物理走線距離，搭配 Hybrid PTFE 壓制 400G SerDes 的高頻損耗，全機櫃內部不進行光電轉換。
- TPFE
 - Switch tray, backplane 都會用到 TPFE ?
- 需求量
 - 一年>1000萬片
 - 台虹每個月25萬片, 每片350
- 2029 貢獻
 - 月營收: 28E
 - 月EPS : 4 元

DIN Kyber TP > 1000, 270% Up,把握:低


定額

生產時間表與驗證狀態矩陣

	驗證/送樣 (3Q26)	量產啟動 (4Q26)	交貨認列 (1Q27)
ASIC 機櫃主板 (20層以上)	已通過驗證 (9M26)	正式轉量產	12M26 出貨
ASIC 高多層板 (30層以上)	量產版本送樣中	預計 12M26 量產	預估 2M27 交貨
美系 GPU 多層板/高階HDI	等待最終回覆	目標 12M26 量產*	

* 時程延後主因客戶量產前多次變更版本，須依新版本重新送驗。
註：高階 AI 板生產週期 (Lead time) 約 6-8 週。

泰國廠毛利率軌跡：量產 S 曲線與目標對齊



現況 (8M26)


稼動率約 60%，月營收 4.3-4.5 億。受惠網通與漲價效益。

量產爬坡 (10-12M26)

高多層訂單進入，營收逐步提升，收斂前期爬坡虧損。

滿載反轉 (1M27)

月產值達 25 億元達成滿載，正式邁向損益兩平以上。



產能開出：

- 新增產能預計於 4Q27 陸續開出，全面因應 GPU/ASIC 需求。

未來潛力：

- P6 廠區已進入規劃，若啟動將於 420 億基礎上再度突破。

Forecast

	2026	2027	2028
Revenue	300E	600E	800E
EPS	0~-1	10	20

12M TP >200, 150% up

표2 GB200 vs GB300 vs VR200 밸류체인 비교			
	GB200 NVL72	GB300 NVL72	VR200 NVL144
Time	▪ Volume production in 2Q25	▪ Volume production in 2026	▪ Volume production in 2H26
Wafer Process	▪ 4nm	▪ 4nm	▪ 3nm
Substrate Supplier	▪ Ibiden, Unimicron	▪ Ibiden, Unimicron	▪ Ibiden, Unimicron
Compute Tray PCB Structure	▪ 5-12-5 (22L HDI)	▪ 5-12-5 (22L HDI)	▪ 6-14-6 (26L HDI)
Compute Tray PCB Suppliers	▪ Unimicron, VGT	▪ VGT, Unimicron	▪ VGT, Unimicron 定穎
NVSwitch Tray PCB Structure	▪ 6+12+6 HDI (24L) ▪ 22L HLC	▪ 22L PCB	▪ 7+12+7 HDI (26L) / 32L HLC
NVSwitch Tray PCB Suppliers	▪ WUS, TTM, VGT, Unimicron	▪ VGT, WUS, Unimicron	▪ VGT, WUS, Unimicron 定穎
CCL Selection	▪ Compute board: M8+M4 ▪ Switch board: M8+M2	▪ Compute board: M8+M4 ▪ Switch board: M8+M2	▪ Compute board: M8.5+M4 ▪ Switch board: M8 ▪ CPX board: M8.5 ▪ Midplane (44L): M8.5/M9
CCL Suppliers	▪ Compute board: Doosan ▪ Switch board: EMC	▪ Compute board: Doosan ▪ Switch board: EMC, SYTECH	▪ Compute board: Doosan ▪ Switch board: EMC, SYTECH ▪ CPX/Midplane: EMC
Copper Foil	▪ HVLP2/3: Mitsui, Circuit Foil, Furukawa	▪ HVLP3: Mitsui, Circuit Foil, Furukawa	▪ HVLP4: Mitsui, Circuit Foil, Furukawa, Co-tech
Fiber Glass	▪ Compute board Low-Dk2: Nittobo, Asahi ▪ Switch board Low-Dk1: Nittobo, Asahi, TWG, Taishan Glass, Grace	▪ Compute board Low-Dk2: Nittobo, Asahi ▪ Switch board Low-Dk1: Nittobo, Asahi, TWG, Taishan Glass, Grace	▪ Low-Dk2/Q-Glass: Nittobo, TWG, Taishan Glass, Grace, CPIC
Resin	▪ Sabic, MGC, Asahi, Pharmicell	▪ Sabic, MGC, Asahi, Pharmicell	▪ Sabic, MGC, Asahi, Pharmicell

자료: Prismark, 메리츠증권 리서치센터